

YUHONG LUO

jamesluoyh.github.io

✉ y.luo@rutgers.edu

🌐 linkedin.com/in/yuhong-luo

🐙 github.com/JamesLuoyh

🔍 [Google Scholar](#)

Education

Rutgers University

Ph.D. Student in Computer Science, GPA: 4.0/4.0

New Brunswick, NJ

Sep. 2024–May 2028 (Expected)

University of Massachusetts Amherst

M.S. in Computer Science, GPA: 3.97/4.0

Amherst, MA

Sep. 2022–May 2024

University of Washington

B.S. in Computer Science and Applied & Computational Mathematical Sciences, GPA: 3.82/4.0

Seattle, WA

Sep. 2016–Mar. 2019

Research Interests

Strategic and multi-agent AI; mechanism design for AI agents; representation engineering and controllable model behavior.

Publications

• Decentralized Aggregation of LLM Predictions via Wagering Mechanisms.

Yuhong Luo, David M. Pennock, and Xintong Wang.

Conference on Neural Information Processing Systems (NeurIPS), 2026. [arXiv] [code/data]

• Algorithmic collusion at inference time: A meta-game design and evaluation.

Yuhong Luo, Daniel Schoepflin, and Xintong Wang.

International Conference on Autonomous Agents and Multiagent Systems (AAMAS), 2026. [arXiv] [proceedings] [code/data]

• Fair Representation Learning with Controllable High Confidence Guarantees via Adversarial Inference.

Yuhong Luo, Austin Hoag, Xintong Wang, Philip S. Thomas, and Przemyslaw A. Grabowicz.

Conference on Neural Information Processing Systems (NeurIPS), 2025. **Scholar Award**. [arXiv] [code/data]

• Scalable and Efficient Temporal Graph Representation Learning via Forward Recent Sampling.

Yuhong Luo and Pan Li.

Learning on Graphs Conference (LoG), 2024. [arXiv] [code/data]

• Neighborhood-aware Scalable Temporal Network Representation Learning.

Yuhong Luo and Pan Li.

Learning on Graphs Conference (LoG), 2022. **Best Paper Award**. [arXiv] [code/data]

• [Preprint] Can LLM Agents Simulate Dynamic Networks? A Case Study on Email Networks with Phishing Synthesis.

Siqi Miao, Ziyang Chen, Yuhong Luo, Hans Hao-Hsun Hsu, Mufei Li, Kaiqing Zhang, and Pan Li.

In submission, 2026. [arXiv]

Ongoing Projects

- **Representation engineering for distributional control:** Characterizing when activation interventions can recover or induce target output distributions, and developing efficient distribution-matching steering methods.
- **Mechanism design for multi-agent LLM planning:** Designing mechanisms for effective planner–executor collaboration under heterogeneous capabilities and private information.

Research Experience

Rutgers University

Graduate Researcher with Prof. Xintong Wang; AI and economics

New Brunswick, NJ

Sep. 2024–Present

University of Massachusetts Amherst

Graduate Researcher with Prof. Philip Thomas; fair representation learning

Amherst, MA

Feb. 2023–Sep. 2023

Graduate Researcher with Prof. Hui Guan; personalized federated learning

Sep. 2022–Dec. 2022

Purdue University & Georgia Tech

Part-time Research Intern with Prof. Pan Li; temporal graph learning

Remote

May 2021–Aug. 2022

University of Washington

Undergraduate Researcher with Prof. Jan Buys; natural language processing

Seattle, WA

Jan. 2019–Mar. 2019

Engineering Experience

Airbnb

Software Engineer, Marketing Technology

San Francisco, CA

Jul. 2019–Sep. 2022

– Developed personalized marketing-content recommendation models and production ML systems serving marketing campaigns.

– Developed internal tools for content authoring, campaign publishing, and coupon management and issuance.

Meta

Software Engineer Intern, ads detection on landing pages

Seattle, WA

Jul. 2018–Sep. 2018

Yahoo

Software Engineer Intern, web services for game alerts

San Francisco, CA

Jul. 2017–Sep. 2017

Invited Talks

- Temporal Graph Reading Group hosted jointly by McGill University, Mila, and the University of Mannheim, 2023. [recording]
- Temporal graph learning, Stanford University, 2022

Awards

- NeurIPS 2025 Scholar Award; LoG 2022 Best Paper Award.

Professional Service

- Conference reviewer: EC 2025/2026; ACL Rolling Review 2026; NeurIPS 2026; ICLR 2027.

Workshops

- **Decentralized Aggregation of LLM Predictions via Wagering Mechanisms.**
Yuhong Luo, David M. Pennock, and Xintong Wang.
ICML: AI Forecasting Workshop, 2026. **Spotlight**; also presented at EC: Game Theory and Mechanism Design with Large Language Models and Marketplace Innovation Workshop (MIW). [arXiv] [code/data]
- **Algorithmic collusion at inference time: A meta-game design and evaluation.**
Yuhong Luo, Daniel Schoepflin, and Xintong Wang.
EC: Game Theory and Mechanism Design with Large Language Models; Marketplace Innovation Workshop (MIW), 2026. [arXiv]